

PostgreSQL Syntax: A Concise but Complete Practical Introduction

(for study use)

September 28, 2025

Contents

1	Mindset & Scope	2
2	Data Types (core + PG extras)	2
3	Schemas, Databases, Search Path	2
4	DDL: Tables, Constraints, Identity, Generated Columns	3
4.1	Create tables	3
4.2	Altering & dropping	3
4.3	Indexes (btree default) and advanced types	3
5	DML: INSERT, UPDATE, DELETE, MERGE, UPSERT	3
6	DQL: SELECT Core	4
6.1	Projection, filtering, ordering, limits	4
6.2	Joins	4
6.3	Aggregates & GROUP BY	4
6.4	Window functions (analytic)	4
6.5	Subqueries & CTEs	5
6.6	Recursive CTE	5
7	Arrays & JSON (Postgres superpowers)	5
7.1	Arrays	5
7.2	JSONB	5
8	Range Types & Exclusion Constraints	5
9	Full-Text Search (FTS)	6
10	Views & Materialized Views	6
11	Sequences & Identity	6
12	Constraints: DEFERRABLE & Triggers	6
13	Partitioning (range, list, hash)	6
14	Transactions (TCL), Isolation, Locks	7
14.1	Basics	7
14.2	Isolation levels	7
14.3	Explicit locks (use sparingly)	7
15	Security (DCL): Roles, Privileges, Row-Level Security	7
16	Query Planning & Performance Hygiene	7

17 Common Table Patterns & Idioms	8
17.1 Upsert with conflict target expression	8
17.2 Keyset pagination	8
17.3 Pivot-ish aggregates	8
17.4 Distinct-on (PG feature)	8
18 Advanced Joins & Set Operations	8
19 Error Handling & Constraints in Practice	8
20 Triggers & PL/pgSQL (sketch)	8
21 Extensions & FDW	9
22 Temporal Patterns	9
23 Admin Odds & Ends (psql meta, safe defaults)	9
24 Common Pitfalls (Nitpicks)	9
25 Minimal End-to-End Example	9
26 Where this fits in a curriculum	10
27 Further Study Pointers (conceptual)	10

1 Mindset & Scope

This guide targets the high-value subset of PostgreSQL syntax you will actually use. It groups syntax by responsibility:

- **DDL** (define structure): CREATE/ALTER/DROP, types, constraints, indexes, partitioning, schemas.
- **DML** (change data): INSERT/UPDATE/DELETE/MERGE.
- **DQL** (query data): SELECT, joins, subqueries, CTE, window functions, aggregates.
- **TCL** (transactions): BEGIN/COMMIT/ROLLBACK, savepoints, isolation.
- **DCL** (control): roles, privileges, GRANT/REVOKE.

Postgres-specifics: **arrays**, **JSON**, **range types**, **full-text search**, **generated columns**, **identity/sequence**, **UPSERT**, **recursive CTEs**, **indexes** (btree/hash/brin/gin/gist), **partitioning**, **locks**, **PL/pgSQL**, **extensions**.

Caveat: “Complete” SQL is unbounded (extensions, planner hints, FDWs). This is the core, with idiomatic examples and pitfalls.

2 Data Types (core + PG extras)

- Numeric: smallint, int, bigint, decimal(p,s), numeric, real, double precision, serial/bigserial (legacy).
- Text/byte: text, varchar(n), char(n), bytea.
- Temporal: date, time, timestamp [with(out) time zone], interval.
- Boolean: boolean.
- UUID, json, jsonb, xml, cidr, inet, macaddr, money.
- Arrays: int[], text[], multidimensional int[][].
- Ranges: int4range, int8range, numrange, tsrange, tstzrange, daterange.
- Composite & enum: CREATE TYPE.

3 Schemas, Databases, Search Path

```
-- create and use a schema
CREATE SCHEMA app AUTHORIZATION app_user;
SET search_path TO app, public; -- resolves unqualified names
```

Pitfall: if you rely on public, lock down CREATE privileges for security.

4 DDL: Tables, Constraints, Identity, Generated Columns

4.1 Create tables

```
CREATE TABLE app.users (  
  user_id bigserial PRIMARY KEY, -- legacy; prefer identity  
  email text NOT NULL UNIQUE,  
  full_name text,  
  is_active boolean NOT NULL DEFAULT true,  
  created_at timestamptz NOT NULL DEFAULT now(),  
  -- identity (SQL standard)  
  v2_id bigint GENERATED ALWAYS AS IDENTITY,  
  -- generated column (stored)  
  email_domain text GENERATED ALWAYS AS (split_part(email, '@', 2)) STORED,  
  -- FK with actions  
  org_id bigint REFERENCES app.orgs(org_id)  
    ON UPDATE CASCADE ON DELETE SET NULL,  
  -- check constraint  
  CONSTRAINT email_ok CHECK (position('@' IN email) > 1)  
);
```

Notes: prefer IDENTITY over serial; generated columns require STORED.

4.2 Altering & dropping

```
ALTER TABLE app.users ADD COLUMN tags text[] DEFAULT '{}':text[];  
ALTER TABLE app.users ALTER COLUMN full_name SET NOT NULL;  
ALTER TABLE app.users DROP COLUMN tags;  
DROP TABLE IF EXISTS app.users CASCADE;
```

4.3 Indexes (btree default) and advanced types

```
-- btree for equality/range  
CREATE INDEX ON app.users (email);  
-- expression index: matches WHERE lower(email)=...  
CREATE INDEX users_email_lower_idx ON app.users ((lower(email)));  
-- partial index: sparse filtered  
CREATE INDEX users_active_idx ON app.users (created_at) WHERE is_active;  
-- unique multi-column  
CREATE UNIQUE INDEX users_org_email_uniq ON app.users (org_id, email);  
  
-- GIN for arrays|jsonb|full-text  
CREATE INDEX users_tags_gin ON app.users USING GIN (tags);  
-- JSONB path ops  
CREATE INDEX users_profile_gin ON app.users USING GIN ((profile jsonb_path_ops));  
-- GiST example (ranges, geo)  
CREATE INDEX events_time_gist ON app.events USING GIST (when_range);  
-- BRIN for very large, naturally clustered tables  
CREATE INDEX logs_brin ON app.logs USING BRIN (ts);
```

Pitfall: LIKE 'prefix%' benefits from btree only with text_pattern_ops or citext / expression indexes (e.g., lower()).

5 DML: INSERT, UPDATE, DELETE, MERGE, UPSERT

```
-- INSERT  
INSERT INTO app.users (email, full_name) VALUES  
  ('a@x.com', 'A'), ('b@x.com', 'B')  
RETURNING user_id;  
  
-- INSERT .. SELECT
```

```

INSERT INTO app.audit(email, at) SELECT email, now() FROM app.users WHERE is_active;

-- UPSERT
INSERT INTO app.users (email, full_name)
VALUES ('a@x.com', 'New Name')
ON CONFLICT (email) DO UPDATE
  SET full_name = EXCLUDED.full_name, updated_at = now();

-- MERGE (PG15+)
MERGE INTO app.inventory AS t
USING (VALUES (1, 'USB-C', 10), (2, 'HDMI', 5)) AS s(id, name, qty)
ON t.id = s.id
WHEN MATCHED THEN UPDATE SET qty = t.qty + s.qty
WHEN NOT MATCHED THEN INSERT (id, name, qty) VALUES (s.id, s.name, s.qty);

-- UPDATE
UPDATE app.users SET is_active = false WHERE last_login < now() - interval '180 days'
RETURNING user_id;

-- DELETE
DELETE FROM app.users WHERE is_active = false AND created_at < now() - interval '1 year';

```

6 DQL: SELECT Core

6.1 Projection, filtering, ordering, limits

```

SELECT user_id, email, created_at
FROM app.users
WHERE is_active AND created_at >= date_trunc('month', now())
ORDER BY created_at DESC
LIMIT 20 OFFSET 0; -- prefer keyset pagination for scale

```

6.2 Joins

```

SELECT u.user_id, u.email, o.name AS org
FROM app.users u
JOIN app.orgs o ON o.org_id = u.org_id
LEFT JOIN app.subscriptions s ON s.user_id = u.user_id
WHERE s.active IS DISTINCT FROM false;

```

Nits: use IS DISTINCT FROM to handle NULL-safe equality.

6.3 Aggregates & GROUP BY

```

SELECT org_id, count(*) AS users, avg(age) AS avg_age
FROM app.users
GROUP BY org_id
HAVING count(*) > 10
ORDER BY users DESC;

```

6.4 Window functions (analytic)

```

SELECT
  user_id, email,
  row_number() OVER (PARTITION BY org_id ORDER BY created_at) AS rn,
  avg(age) OVER (PARTITION BY org_id) AS org_avg_age
FROM app.users;

```

Concept link: aggregation collapses rows; windows compute per-row analytics without collapsing.

6.5 Subqueries & CTEs

```
WITH active AS (  
    SELECT user_id, org_id FROM app.users WHERE is_active  
)  
SELECT o.name, count(a.user_id)  
FROM active a JOIN app.orgs o USING(org_id)  
GROUP BY o.name;
```

Nitpick: CTEs are *optimization fences* before PG12; from PG12 they can inline. Avoid unnecessary CTEs in hot paths.

6.6 Recursive CTE

```
WITH RECURSIVE org_tree AS (  
    SELECT org_id, parent_id, name, 1 AS lvl  
    FROM app.orgs WHERE parent_id IS NULL  
    UNION ALL  
    SELECT c.org_id, c.parent_id, c.name, p.lvl+1  
    FROM app.orgs c JOIN org_tree p ON c.parent_id = p.org_id  
)  
SELECT * FROM org_tree ORDER BY lvl, name;
```

7 Arrays & JSON (Postgres superpowers)

7.1 Arrays

```
-- literals and operators  
SELECT ARRAY[1,2,3] @> ARRAY[2]; -- contains  
SELECT 2 = ANY(ARRAY[1,2,3]); -- membership  
-- unnest with ordinality for position-aware explode  
SELECT email, tag, ord  
FROM app.users u,  
LATERAL unnest(u.tags) WITH ORDINALITY AS t(tag, ord);
```

Indexing: GIN on arrays supports @>, &&, etc.

7.2 JSONB

```
-- access  
SELECT profile->>'lang' AS lang  
FROM app.users  
WHERE profile @> '{"marketing": {"opt_in": true}}';  
  
-- jsonpath (PG12+)  
SELECT jsonb_path_query_array(profile, '$.addresses[*].city')  
FROM app.users;  
  
-- GIN indexing for containment queries  
CREATE INDEX users_profile_gin ON app.users USING GIN (profile jsonb_path_ops);
```

Tradeoff: prefer normalized schema for frequently filtered fields; use JSONB for sparse or evolving attributes.

8 Range Types & Exclusion Constraints

```
CREATE TABLE room_bookings(  
    room_id int, during tstzrange NOT NULL,  
    EXCLUDE USING GIST (room_id WITH =, during WITH &&) -- disallow overlaps  
);  
-- include/exclude bounds
```

```
INSERT INTO room_bookings VALUES (101, tstzrange(now(), now()+interval '2h','[]'));
```

9 Full-Text Search (FTS)

```
ALTER TABLE docs ADD COLUMN tsv tsvector;  
UPDATE docs SET tsv = to_tsvector('english', coalesce(title,'') || ' ' || coalesce(body,''));  
CREATE INDEX docs_tsv_idx ON docs USING GIN (tsv);  
SELECT * FROM docs WHERE tsv @@ plainto_tsquery('english', 'quick fox');
```

Note: maintain tsv via trigger for write paths.

10 Views & Materialized Views

```
CREATE VIEW app.active_users AS  
  SELECT * FROM app.users WHERE is_active;  
  
CREATE MATERIALIZED VIEW app.monthly_summary AS  
  SELECT date_trunc('month', created_at) AS m, count(*) c  
  FROM app.users GROUP BY 1;  
  
REFRESH MATERIALIZED VIEW CONCURRENTLY app.monthly_summary; -- needs unique index
```

11 Sequences & Identity

```
CREATE SEQUENCE app.seq START 100 INCREMENT 10 OWNED BY app.users.user_id;  
SELECT nextval('app.seq');  
  
-- identity columns (preferred)  
CREATE TABLE t(i int GENERATED BY DEFAULT AS IDENTITY);
```

12 Constraints: DEFERRABLE & Triggers

```
ALTER TABLE app.users  
  ADD CONSTRAINT org_fk FOREIGN KEY (org_id) REFERENCES app.orgs(org_id)  
  DEFERRABLE INITIALLY DEFERRED;  
  
BEGIN;  
SET CONSTRAINTS org_fk DEFERRED;  
-- batch of rows becoming valid only at txn end  
COMMIT;
```

13 Partitioning (range, list, hash)

```
CREATE TABLE events(  
  id bigserial, ts timestamptz NOT NULL, payload jsonb  
) PARTITION BY RANGE (ts);  
  
CREATE TABLE events_2025_09 PARTITION OF events  
  FOR VALUES FROM ('2025-09-01') TO ('2025-10-01');  
  
-- indexing on parent propagates to partitions in PG11+  
CREATE INDEX ON events (ts);
```

Use-case: time-series/logs; enables pruning and lighter maintenance.

14 Transactions (TCL), Isolation, Locks

14.1 Basics

```
BEGIN;
-- ... DML ...
SAVEPOINT s1;
-- ... attempt risky step ...
ROLLBACK TO SAVEPOINT s1;
COMMIT;
```

14.2 Isolation levels

```
SET TRANSACTION ISOLATION LEVEL READ COMMITTED; -- default
-- REPEATABLE READ (no non-repeatable reads; phantom-free in PG)
-- SERIALIZABLE (true serializability; may abort on conflicts)
```

Pitfall: SERIALIZABLE can abort; retry logic required.

14.3 Explicit locks (use sparingly)

```
SELECT * FROM app.users WHERE user_id=42 FOR UPDATE; -- row lock
LOCK TABLE app.users IN SHARE MODE; -- table lock
```

15 Security (DCL): Roles, Privileges, Row-Level Security

```
-- roles
CREATE ROLE app_user LOGIN PASSWORD '...';
GRANT USAGE ON SCHEMA app TO app_user;
GRANT SELECT, INSERT, UPDATE ON ALL TABLES IN SCHEMA app TO app_user;

-- future-proof default privileges
ALTER DEFAULT PRIVILEGES IN SCHEMA app
    GRANT SELECT, INSERT, UPDATE ON TABLES TO app_user;

-- Row-Level Security
ALTER TABLE app.users ENABLE ROW LEVEL SECURITY;
CREATE POLICY same_org ON app.users
USING (org_id = current_setting('app.current_org')::int);
```

Link: RLS complements application auth; be explicit about ALTER TABLE ... FORCE ROW LEVEL SECURITY.

16 Query Planning & Performance Hygiene

```
EXPLAIN (ANALYZE, BUFFERS, VERBOSE)
SELECT ... ;

-- maintenance
VACUUM (ANALYZE); -- autovacuum usually handles this
SET work_mem='64MB'; -- per-sort/hash; beware overcommit
SET enable_seqscan=off; -- debugging only (do not use in prod)
```

Mental model: Pick *right data types, right indexes, right predicates*. Prefer keyset pagination over OFFSET for large pages.

17 Common Table Patterns & Idioms

17.1 Upsert with conflict target expression

```
CREATE UNIQUE INDEX users_email_lower ON app.users ((lower(email)));
INSERT INTO app.users(email) VALUES('A@X.com')
ON CONFLICT ((lower(email))) DO NOTHING;
```

17.2 Keyset pagination

```
SELECT * FROM app.users
WHERE created_at < $last_seen
ORDER BY created_at DESC
LIMIT 50;
```

17.3 Pivot-ish aggregates

```
SELECT org_id,
       count(*) FILTER (WHERE is_active) AS active,
       count(*) FILTER (WHERE NOT is_active) AS inactive
FROM app.users GROUP BY org_id;
```

17.4 Distinct-on (PG feature)

```
SELECT DISTINCT ON (org_id) org_id, user_id, created_at
FROM app.users
ORDER BY org_id, created_at DESC;
```

18 Advanced Joins & Set Operations

```
-- set ops: UNION [ALL], INTERSECT, EXCEPT
SELECT email FROM a
UNION
SELECT email FROM b;

-- lateral joins: use outputs of a subquery per-row
SELECT u.user_id, f.friend_id
FROM app.users u
LEFT JOIN LATERAL (
  SELECT friend_id FROM app.friends WHERE user_id=u.user_id ORDER BY since DESC LIMIT 5
) f ON true;
```

19 Error Handling & Constraints in Practice

- Prefer NOT NULL + defaults; CHECK for domain rules.
- Use DEFERRABLE when intra-batch temporary violations are expected.
- Validate emails/phones at app layer; use light DB checks only.

20 Triggers & PL/pgSQL (sketch)


```
CREATE OR REPLACE FUNCTION app.users_tsv_update()
RETURNS trigger LANGUAGE plpgsql AS $$
BEGIN
    NEW.tsv := to_tsvector('english', coalesce(NEW.full_name, '') || ' ' || coalesce(NEW.email, ''));
    RETURN NEW;
END $$;

CREATE TRIGGER users_tsv_trg
BEFORE INSERT OR UPDATE OF full_name, email ON app.users
FOR EACH ROW EXECUTE FUNCTION app.users_tsv_update();
```

Rule of thumb: keep business logic in app layer unless DB-centric.

21 Extensions & FDW

```
CREATE EXTENSION IF NOT EXISTS "uuid-ossf";
CREATE EXTENSION IF NOT EXISTS pg_trgm; -- trigram search
CREATE EXTENSION IF NOT EXISTS postgres_fdw;
```

22 Temporal Patterns

```
-- system-versioning pattern (manual):
valid_from timestamptz NOT NULL DEFAULT now(),
valid_to timestamptz,
EXCLUDE USING GIST (id WITH =, tstzrange(valid_from, coalesce(valid_to, 'infinity')) WITH &&)
```

23 Admin Odds & Ends (psql meta, safe defaults)

```
-- psql (client) meta-commands (not SQL):
-- \d app.users describe
-- \dt app.* list tables
-- \di list indexes
-- \df list functions
```

Tip: in teaching labs, align with NUS/CMU-style DB courses: write EXPLAIN ANALYZE for each nontrivial query and justify index choice.

24 Common Pitfalls (Nitpicks)

- **NULL traps:** = vs IS NULL; use IS DISTINCT FROM.
- **Text & collation:** cross-locale comparisons may disable index usage; normalize with lower() + expression index or citext.
- **LIKE indexing:** %*pattern* (leading %) cannot use btree; consider trigram (pg_trgm).
- **Offset pagination:** avoid large OFFSET; use keyset.
- **CTE misuse:** avoid as “style”; prefer inline subqueries for performance unless readability & reuse dominate.
- **Serial vs identity:** prefer GENERATED {ALWAYS|BY DEFAULT} AS IDENTITY.
- **Transactions:** remember that EXCEPTION in PL/pgSQL subtransaction rolls back only to block.

25 Minimal End-to-End Example

```
BEGIN;
CREATE SCHEMA app;
CREATE TABLE app.orgs (
    org_id bigserial PRIMARY KEY,
    name text NOT NULL UNIQUE
```

```

);
CREATE TABLE app.users (
  user_id bigint GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
  org_id bigint NOT NULL REFERENCES app.orgs(org_id),
  email text NOT NULL,
  created_at timestamptz NOT NULL DEFAULT now(),
  CONSTRAINT email_uniq_per_org UNIQUE (org_id, email)
);
CREATE INDEX ON app.users (created_at);
INSERT INTO app.orgs(name) VALUES ('NUS'),('KTH') RETURNING org_id;

WITH ins AS (
  INSERT INTO app.users (org_id, email)
  SELECT org_id, concat('user', g, '@x.com')
  FROM app.orgs, generate_series(1,3) AS g
  RETURNING *
)
SELECT * FROM ins;

-- query: last 2 users per org
SELECT * FROM (
  SELECT org_id, email, created_at,
         row_number() OVER (PARTITION BY org_id ORDER BY created_at DESC) AS rn
  FROM app.users
) t WHERE rn <= 2;

COMMIT;

```

26 Where this fits in a curriculum

This mirrors typical university DB courses (e.g., Berkeley CS186, CMU 15-415, NUS CS2102): master relational algebra in practice (SELECT/FROM/WHERE/GROUP BY/HAVING), then PG-specific power features (arrays/JSON/windows/CTEs/indexes), and finally transactional semantics (isolation/locking/RLS).

27 Further Study Pointers (conceptual)

Normalization vs. JSONB tradeoffs; logical vs. physical design; cost-based planning; index selectivity; MVCC internals; vacuum and freeze; partition pruning; EXPLAIN reading; concurrency anomalies & serializable retries; FDW performance; pg_trgm/btree_gin/btree_gist nuances.